

# book-rental-recommendation

March 24, 2023

## Project 3 - Book Rental Recommendation

### DESCRIPTION

Book Rent is the largest online and offline book rental chain in India. They provide books of various genres, such as thrillers, mysteries, romances, and science fiction. The company charges a fixed rental fee for a book per month. Lately, the company has been losing its user base. The main reason for this is that users are not able to choose the right books for themselves. The company wants to solve this problem and increase its revenue and profit.

### Project Objective:

You, as an ML expert, should focus on improving the user experience by personalizing it to the user's needs. You have to model a recommendation engine so that users get recommendations for books based on the behavior of similar users. This will ensure that users are renting the books based on their tastes and traits.

Note: You have to perform user-based collaborative filtering and item-based collaborative filtering.

Dataset Description: \* BX-Users: It contains the information of users. \* user\_id - These have been anonymized and mapped to integers \* Location - Demographic data is provided \* Age - Demographic data is provided If available. Otherwise, these fields contain NULL-values.

- BX-Books:
  - isbn - Books are identified by their respective ISBNs. Invalid ISBNs have already been removed from the dataset.
  - book\_title
  - book\_author
  - year\_of\_publication
  - publisher
- BX-Book-Ratings: Contains the book rating information.
  - user\_id
  - isbn
  - rating - Ratings (**Book-Rating**) are either explicit, expressed on a scale from 1-10 (higher values denoting higher appreciation), or implicit, expressed by 0.

Note: Download the “BX-Book-Ratings.csv”, “BX-Books.csv”, “BX-Users.csv”, and “Recommend.csv” using the link given in the Book Rental Recommendation project problem statement.

Following operations should be performed:

- Read the books dataset and explore it
- Clean up NaN values

- Read the data where ratings are given by users
- Take a quick look at the number of unique users and books
- Convert ISBN variables to numeric numbers in the correct order
- Convert the user\_id variable to numeric numbers in the correct order
- Convert both user\_id and ISBN to the ordered list, i.e., from 0...n-1
- Re-index the columns to build a matrix
- Split your data into two sets (training and testing)
- Make predictions based on user and item variables
- Use RMSE to evaluate the predictions

```
[1]: import numpy as np
import pandas as pd

import warnings
warnings.filterwarnings('ignore')
```

Read the data using pandas in DataFrame df\_user To map byte values directly to the first 256 Unicode code points, use the “Latin-1” encoding. This is the closest equivalent Python 3 offers to the permissive Python 2 text handling model.

```
[2]: df_user = pd.read_csv('BX-Users.csv',encoding='latin-1')
```

```
[3]: df_user.head()
```

```
[3]:   user_id      Location  Age
0        1  nyc, new york, usa  NaN
1        2  stockton, california, usa  18.0
2        3  moscow, yukon territory, russia  NaN
3        4  porto, v.n.gaia, portugal  17.0
4        5  farnborough, hants, united kingdom  NaN
```

```
[4]: df_user.tail()
```

```
[4]:   user_id      Location  Age
278854  278854  portland, oregon, usa  NaN
278855  278855  tacoma, washington, united kingdom  50.0
278856  278856  brampton, ontario, canada  NaN
278857  278857  knoxville, tennessee, usa  NaN
278858  278858  dublin, n/a, ireland  NaN
```

```
[5]: df_user.shape
```

```
[5]: (278859, 3)
```

```
[6]: df_user.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 278859 entries, 0 to 278858
```

```
Data columns (total 3 columns):
#   Column      Non-Null Count  Dtype
---  -
0   user_id     278859 non-null   object
1   Location     278858 non-null   object
2   Age          168096 non-null   float64
dtypes: float64(1), object(2)
memory usage: 6.4+ MB
```

```
[7]: # Checking for Null Values.
df_user.isnull().sum()
```

```
[7]: user_id      0
Location      1
Age          110763
dtype: int64
```

```
[8]: df_user.isnull().any()
```

```
[8]: user_id      False
Location      True
Age           True
dtype: bool
```

```
[9]: # Dropping the Null Values.
df_user1=df_user.dropna()
```

```
[10]: df_user1.isnull().sum()
```

```
[10]: user_id      0
Location      0
Age           0
dtype: int64
```

```
[11]: df_user1.isnull().any()
```

```
[11]: user_id      False
Location      False
Age           False
dtype: bool
```

Read the books Data and explore

```
[12]: df_books = pd.read_csv('BX-Books.csv', encoding='latin-1')
```

```
[13]: df_books.head()
```

```
[13]:          isbn                                book_title \
0  195153448                                Classical Mythology
1    2005018                                Clara Callan
2   60973129                                Decision in Normandy
3  374157065  Flu: The Story of the Great Influenza Pandemic...
4  393045218                                The Mummies of Urumchi

          book_author year_of_publication publisher
0    Mark P. O. Morford            2002  Oxford University Press
1  Richard Bruce Wright            2001   HarperFlamingo Canada
2         Carlo D'Este            1991     HarperPerennial
3    Gina Bari Kolata            1999   Farrar Straus Giroux
4     E. J. W. Barber            1999  W. W. Norton & Company
```

```
[14]: df_books.shape
```

```
[14]: (271379, 5)
```

Reading the data where ratings are given *We will read only first 10000 rows otherwise, Out Of Memory error can occur.*

```
[15]: df_ratings = pd.read_csv('BX-Book-Ratings.csv',encoding='latin-1',nrows=10000)
```

```
[16]: df_ratings.head()
```

```
[16]:   user_id      isbn  rating
0   276725  034545104X      0
1   276726  155061224      5
2   276727  446520802      0
3   276729  052165615X      3
4   276729  521795028      6
```

Using ‘**describe()**’ function *It is used to view some basic statistical details like percentile, mean, std.*

```
[17]: df_ratings.describe()
```

```
[17]:
```

	user_id	rating
count	10000.000000	10000.000000
mean	265844.379600	1.974700
std	56937.189618	3.424884
min	2.000000	0.000000
25%	277478.000000	0.000000
50%	278418.000000	0.000000
75%	278418.000000	4.000000
max	278854.000000	10.000000

Merge the dataframes *For all practical purposes, User Master Data is not required. So, ignore*

*dataframe df\_user*

```
[18]: df_final = pd.merge(df_ratings,df_books,on='isbn')
```

```
[19]: df_final.head()
```

```
[19]:
```

	user_id	isbn	rating	book_title	book_author \
0	276725	034545104X	0	Flesh Tones: A Novel	M. J. Rose
1	276726	155061224	5	Rites of Passage	Judith Rae
2	276727	446520802	0	The Notebook	Nicholas Sparks
3	278418	446520802	0	The Notebook	Nicholas Sparks
4	276729	052165615X	3	Help!: Level 1	Philip Prowse

  

	year_of_publication	publisher
0	2002	Ballantine Books
1	2001	Heinle
2	1996	Warner Books
3	1996	Warner Books
4	1999	Cambridge University Press

Checking for unique users and books Here we are using ‘**nunique()**’ function that returns the Series with the number of distinct observations over the requested axis.

```
[20]: # Code for checking number of unique users and books.
n_users = df_final.user_id.nunique()
n_books = df_final.isbn.nunique()

print('Num. of Users: '+ str(n_users))
print('Num of Books: '+str(n_books))
```

Num. of Users: 828

Num of Books: 8051

Convert ISBN variable to numeric type in order

```
[21]: # Convert and print length of isbn list
isbn_list = df_final.isbn.unique()
print(" Length of isbn List:", len(isbn_list))
def get_isbn_numeric_id(isbn):
    #print (" isbn is:" , isbn)
    itemindex = np.where(isbn_list==isbn)
    return itemindex[0][0]
```

Length of isbn List: 8051

Convert user\_id variable to numeric type in order *This is formatted as code.*

```
[22]: # Convert and print length of user_id list
userid_list = df_final.user_id.unique()
```

```

print(" Length of user_id List:", len(userid_list))
def get_user_id_numeric_id(user_id):
    #print (" isbn is:" , isbn)
    itemindex = np.where(userid_list==user_id)
    return itemindex[0][0]

```

Length of user\_id List: 828

Convert both user\_id and isbn to ordered list i.e. from 0..n-1

```
[23]: df_final['user_id_order'] = df_final['user_id'].apply(get_user_id_numeric_id)
```

```
[24]: df_final['isbn_id'] = df_final['isbn'].apply(get_isbn_numeric_id)
df_final.head()
```

```
[24]:
```

	user_id	isbn	rating	book_title	book_author \
0	276725	034545104X	0	Flesh Tones: A Novel	M. J. Rose
1	276726	155061224	5	Rites of Passage	Judith Rae
2	276727	446520802	0	The Notebook	Nicholas Sparks
3	278418	446520802	0	The Notebook	Nicholas Sparks
4	276729	052165615X	3	Help!: Level 1	Philip Prowse

  

	year_of_publication		publisher	user_id_order	isbn_id
0		2002	Ballantine Books	0	0
1		2001	Heinle	1	1
2		1996	Warner Books	2	2
3		1996	Warner Books	3	2
4		1999	Cambridge University Press	4	3

Re-index columns to build matrix

```
[25]: # Reindexing the columns
new_col_order = ['user_id_order', 'isbn_id', 'rating', 'book_title', 'book_author', 'year_of_publication', 'publisher', 'isbn', 'user_id']
df_final = df_final.reindex(columns= new_col_order)
df_final.head()
```

```
[25]:
```

	user_id_order	isbn_id	rating	book_title	book_author \
0	0	0	0	Flesh Tones: A Novel	M. J. Rose
1	1	1	5	Rites of Passage	Judith Rae
2	2	2	0	The Notebook	Nicholas Sparks
3	3	2	0	The Notebook	Nicholas Sparks
4	4	3	3	Help!: Level 1	Philip Prowse

  

	year_of_publication		publisher	isbn	user_id
0		2002	Ballantine Books	034545104X	276725
1		2001	Heinle	155061224	276726
2		1996	Warner Books	446520802	276727

3	1996	Warner Books	446520802	278418
4	1999	Cambridge University Press	052165615X	276729

## Train Test Split

Recommendation Systems are difficult to evaluate, but we will still learn how to evaluate them. In order to do this, will split our data into two sets. However, we won't do our classic `X_train, X_test, y_train, y_test` split. Instead, we can actually just segment the data into two sets of data:

Importing `train_test_split` model

```
[26]: # Importing train_test_split model for splittig the data into train and test_
      ↪ set.
      from sklearn.model_selection import train_test_split
      train_data, test_data = train_test_split(df_final, test_size=0.20)
```

Approach: We Will Use Memory-Based Collaborative Filtering

Memory-Based Collaborative Filtering approaches can be divided into two main sections: **user-item filtering** and **item-item filtering**.

A *user-item filtering* will take a particular user, find users that are similar to that user based on similarity of ratings, and recommend items that those similar users liked.

In contrast, *item-item filtering* will take an item, find users who liked that item, and find other items that those users or similar users also liked. It takes items as input and outputs other items as recommendations.

- *Item-Item Collaborative Filtering*: “Users who liked this item also liked ...”
- *User-Item Collaborative Filtering*: “Users who are similar to you also liked ...”

In both cases, we will create a user-book matrix which is built from the entire dataset. Since we have split the data into testing and training, we will need to create two [828 x 8051] matrices (all users by all books). This is going to be a very large matrix. The training matrix contains 80% of the ratings and the testing matrix contains 20% of the ratings.

Create two user-book matrix for training and testing

Indented block

```
[27]: # Create user-book matrix for training
      train_data_matrix = np.zeros((n_users, n_books))
      for line in train_data.itertuples():
          train_data_matrix[line[1]-1, line[2]-1] = line[3]

      # Create user-book matrix for testing
      test_data_matrix = np.zeros((n_users, n_books))
      for line in test_data.itertuples():
          test_data_matrix[line[1]-1, line[2]-1] = line[3]
```

Import Pairwise Model we can use the `pairwise_distances` function from `sklearn` to calculate the cosine similarity. Note, the output will range from 0 to 1 since the ratings are all positive.

```
[28]: # Importing pairwise_distances function
from sklearn.metrics.pairwise import pairwise_distances
user_similarity = pairwise_distances(train_data_matrix, metric='cosine')
item_similarity = pairwise_distances(train_data_matrix.T, metric='cosine')
```

```
[29]: user_similarity
```

```
[29]: array([[0., 1., 1., ..., 1., 1., 1.],
        [1., 0., 1., ..., 1., 1., 1.],
        [1., 1., 0., ..., 1., 1., 1.],
        ...,
        [1., 1., 1., ..., 0., 1., 1.],
        [1., 1., 1., ..., 1., 0., 1.],
        [1., 1., 1., ..., 1., 1., 0.]])
```

Make predictions

```
[30]: # Defining custom function to make predictions
def predict(ratings, similarity, type='user'):
    if type == 'user':
        mean_user_rating = ratings.mean(axis=1)
        # We will use np.newaxis so that mean_user_rating has same format as
        ratings.
        ratings_diff = (ratings - mean_user_rating[:, np.newaxis])
        pred = mean_user_rating[:, np.newaxis] + similarity.dot(ratings_diff) /
        np.array([np.abs(similarity).sum(axis=1)]).T
    elif type == 'item':
        pred = ratings.dot(similarity) / np.array([np.abs(similarity).
        sum(axis=1)])
    return pred
```

```
[31]: item_prediction = predict(train_data_matrix, item_similarity, type='item')
user_prediction = predict(train_data_matrix, user_similarity, type='user')
```

```
[32]: print(item_prediction)
```

```
[[0.          0.00062112 0.00062112 ... 0.00062159 0.00062112 0.00062112]
 [0.          0.          0.          ... 0.          0.          0.          ]
 [0.06571429 0.06571429 0.06571429 ... 0.06576407 0.06571429 0.06571429]
 ...
 [0.          0.          0.          ... 0.          0.          0.          ]
 [0.          0.          0.          ... 0.          0.          0.          ]
 [0.          0.          0.          ... 0.          0.          0.          ]]
```

```
[33]: print(user_prediction)
```

```
[[ -0.00135983 -0.00135983  0.00226774 ...  0.00952288 -0.00135983
  -0.00135983]]
```



```
[ 0.00406433 -0.00198162  0.00164595 ...  0.00890109 -0.00198162
 -0.00198162]
[ 0.06985073  0.06380399  0.06743203 ...  0.07468811  0.06380399
 0.06380399]
...
[ 0.00406433 -0.00198162  0.00164595 ...  0.00890109 -0.00198162
 -0.00198162]
[ 0.00406433 -0.00198162  0.00164595 ...  0.00890109 -0.00198162
 -0.00198162]
[ 0.00406433 -0.00198162  0.00164595 ...  0.00890109 -0.00198162
 -0.00198162]]
```

Evaluation There are many evaluation metrics, but one of the most popular metric used to evaluate accuracy of predicted ratings is *Root Mean Squared Error (RMSE)*.

Since, we only want to consider predicted ratings that are in the test dataset, we will filter out all other elements in the prediction matrix with: `prediction[ground_truth.nonzero()]`.

```
[34]: # Importing RMSE function
from sklearn.metrics import mean_squared_error
from math import sqrt

# Defining custom function to filter out elements with ground_truth.nonzero
def rmse(prediction, ground_truth):
    prediction = prediction[ground_truth.nonzero()].flatten()
    ground_truth = ground_truth[ground_truth.nonzero()].flatten()
    return sqrt(mean_squared_error(prediction, ground_truth))
```

Printing RMSE value for user based and item based collaborative filtering

```
[35]: print('User-based CF RMSE: ' + str(rmse(user_prediction, test_data_matrix)))
print('Item-based CF RMSE: ' + str(rmse(item_prediction, test_data_matrix)))
```

User-based CF RMSE: 7.663374330427166

Item-based CF RMSE: 7.662434131747383

Both the approach yield almost same result.